

Preparing Data And Preprocessing

```

In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.ensemble import RandomForestRegressor
from sklearn.decomposition import PCA
from sklearn.cluster import DBSCAN
from sklearn.preprocessing import StandardScaler, PowerTransformer
from sklearn.feature_selection import VarianceThreshold
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.neighbors import LocalOutlierFactor
from kneed import KneeLocator
from sklearn.cluster import KMeans
from sklearn.mixture import GaussianMixture
import umap.umap_ as umap

# ---- STEP 1: LOAD DATA ----
df = pd.read_excel('./data_files/Data_re.xlsx') # Ensure the file exists

# ---- STEP 2: FEATURE SELECTION ----
cols_to_keep = ['g_flux', 'r_flux', 'i_flux', 'y_flux', 'z_flux', 'specz_
df_selected = df[cols_to_keep].copy()

# Compute Flux Color Indices
df_selected['g_r'] = df_selected['g_flux'] - df_selected['r_flux']
df_selected['r_i'] = df_selected['r_flux'] - df_selected['i_flux']
df_selected['i_y'] = df_selected['i_flux'] - df_selected['y_flux']
df_selected['y_z'] = df_selected['y_flux'] - df_selected['z_flux']

# ---- STEP 3: REMOVE OUTLIERS ----
lof = LocalOutlierFactor(n_neighbors=20, contamination=0.02)
outlier_scores = lof.fit_predict(df_selected.drop(columns=['specz_redshif
df_clean = df_selected[outlier_scores == 1].copy()

# ---- STEP 4: SCALING & TRANSFORMING ----
scaler = StandardScaler()
df_scaled = scaler.fit_transform(df_clean.drop(columns=['specz_redshift'])

# Apply Power Transformation (Yeo-Johnson for normalizing skewed data)
power_transformer = PowerTransformer(method='yeo-johnson')
df_transformed = power_transformer.fit_transform(df_scaled)

# ---- STEP 5: DIMENSIONALITY REDUCTION WITH UMAP ----
reducer = umap.UMAP(n_components=2, n_neighbors=15, min_dist=0.1, random_
X_umap = reducer.fit_transform(df_transformed)

# ---- STEP 6: OPTIONAL PCA FOR EXPLORATION ----
pca = PCA(n_components=3)
X_pca = pca.fit_transform(df_transformed)

# ---- STEP 7: CLUSTERING ----
kmeans = KMeans(n_clusters=3, random_state=42)
df_clean['Cluster'] = kmeans.fit_predict(X_umap)

# ---- PLOT RESULTS ----

```

```

2025-02-21 19:44:36.244662: E external/local_xla/xla/stream_executor/cuda/
cuda_fft.cc:477] Unable to register cuFFT factory: Attempting to register
factory for plugin cuFFT when one has already been registered
WARNING: All log messages before absl::InitializeLog() is called are writt
en to STDERR
E0000 00:00:1740147276.299220      4415 cuda_dnn.cc:8310] Unable to register
cuDNN factory: Attempting to register factory for plugin cuDNN when one ha
s already been registered
E0000 00:00:1740147276.314946      4415 cuda_blas.cc:1418] Unable to registe
r cuBLAS factory: Attempting to register factory for plugin cuBLAS when on
e has already been registered
2025-02-21 19:44:36.430661: I tensorflow/core/platform/cpu_feature_guard.c
c:210] This TensorFlow binary is optimized to use available CPU instructio
ns in performance-critical operations.
To enable the following instructions: AVX2 FMA, in other operations, rebui
ld TensorFlow with the appropriate compiler flags.
/home/chloy/miniconda3/lib/python3.10/site-packages/sklearn/utils/deprecat
ion.py:151: FutureWarning: 'force_all_finite' was renamed to 'ensure_all_f
inite' in 1.6 and will be removed in 1.8.
  warnings.warn(
/home/chloy/miniconda3/lib/python3.10/site-packages/umap/umap_.py:1952: Us
erWarning: n_jobs value 1 overridden to 1 by setting random_state. Use no
seed for parallelism.
  warn(

```



```
In [2]: print(df.head())  
        print(df.info())  
        print(df.describe())
```

```

                                coord      dec  \
0  b'(179325.3125, 99694.8046875, -21178.96484375)' -5.893433
1  b'(179236.609375, 99349.8203125, -23431.181640... -6.522742
2      b'(179281.5, 99283.5078125, -23368.759765625)' -6.505290
3  b'(179365.171875, 99158.046875, -23259.0839843... -6.474628
4      b'(179366.421875, 99172.25, -23188.84765625)' -6.454993

g_central_image_pop_10px_rad  g_central_image_pop_15px_rad  \
0                               1                               1
1                               1                               1
2                               1                               1
3                               1                               1
4                               1                               1

g_central_image_pop_5px_rad  g_cmodel_mag  g_cmodel_magsigma  \
0                               1      20.314907      0.002624
1                               1      22.217360      0.010902
2                               1      21.148739      0.008013
3                               1      18.464205      0.001740
4                               1      20.998287      0.006011

g_ellipticity  g_half_light_radius  g_isophotal_area  ...  z_minor_axis
\
0          0.147                6.047                603  ...      4.938
1          0.130                3.430                93   ...      2.713
2          0.209                6.597                254  ...      4.351
3          0.525               10.855               1064  ...      4.815
4          0.738                8.261                386  ...      2.279

z_peak_surface_brightness  z_petro_rad  z_pos_angle  z_sersic_index  \
0          -8.2933           5.28           36.15           2.193
1          -7.3657           5.94          -61.78           1.649
2          -7.6539           9.24           32.76           2.364
3          -8.5825           6.60           53.15           1.494
4          -6.8798           5.94           21.16           1.063

g_flux      r_flux      i_flux      y_flux      z_flux
0  7.482335e-09  2.987891e-08  5.117533e-08  8.117985e-08  6.737572e-08
1  1.297347e-09  5.643225e-09  1.319162e-08  2.141019e-08  1.828079e-08
2  3.471398e-09  1.571304e-08  3.743521e-08  6.504857e-08  5.585498e-08
3  4.114510e-08  1.090914e-07  1.841788e-07  3.080686e-07  2.402212e-07
4  3.987357e-09  1.373544e-08  2.634882e-08  4.449762e-08  3.527128e-08

```

[5 rows x 89 columns]

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 40914 entries, 0 to 40913

Data columns (total 89 columns):

#	Column	Non-Null Count	Dtype
0	coord	40914 non-null	object
1	dec	40914 non-null	float64
2	g_central_image_pop_10px_rad	40914 non-null	int64
3	g_central_image_pop_15px_rad	40914 non-null	int64
4	g_central_image_pop_5px_rad	40914 non-null	int64
5	g_cmodel_mag	40914 non-null	float64
6	g_cmodel_magsigma	40914 non-null	float64
7	g_ellipticity	40914 non-null	float64
8	g_half_light_radius	40914 non-null	float64
9	g_isophotal_area	40914 non-null	int64
10	g_major_axis	40914 non-null	float64

11	g_minor_axis	40914	non-null	float64
12	g_peak_surface_brightness	40914	non-null	float64
13	g_petro_rad	40914	non-null	float64
14	g_pos_angle	40914	non-null	float64
15	g_sersic_index	40914	non-null	float64
16	i_central_image_pop_10px_rad	40914	non-null	int64
17	i_central_image_pop_15px_rad	40914	non-null	int64
18	i_central_image_pop_5px_rad	40914	non-null	int64
19	i_cmodel_mag	40914	non-null	float64
20	i_cmodel_magsigma	40914	non-null	float64
21	i_ellipticity	40914	non-null	float64
22	i_half_light_radius	40914	non-null	float64
23	i_isophotal_area	40914	non-null	int64
24	i_major_axis	40914	non-null	float64
25	i_minor_axis	40914	non-null	float64
26	i_peak_surface_brightness	40914	non-null	float64
27	i_petro_rad	40914	non-null	float64
28	i_pos_angle	40914	non-null	float64
29	i_sersic_index	40914	non-null	float64
30	object_id	40914	non-null	int64
31	r_central_image_pop_10px_rad	40914	non-null	int64
32	r_central_image_pop_15px_rad	40914	non-null	int64
33	r_central_image_pop_5px_rad	40914	non-null	int64
34	r_cmodel_mag	40914	non-null	float64
35	r_cmodel_magsigma	40914	non-null	float64
36	r_ellipticity	40914	non-null	float64
37	r_half_light_radius	40914	non-null	float64
38	r_isophotal_area	40914	non-null	int64
39	r_major_axis	40914	non-null	float64
40	r_minor_axis	40914	non-null	float64
41	r_peak_surface_brightness	40914	non-null	float64
42	r_petro_rad	40914	non-null	float64
43	r_pos_angle	40914	non-null	float64
44	r_sersic_index	40914	non-null	float64
45	ra	40914	non-null	float64
46	skymap_id	40914	non-null	int64
47	specz_dec	40914	non-null	float64
48	specz_flag_homogeneous	40914	non-null	bool
49	specz_mag_i	40914	non-null	float64
50	specz_name	40914	non-null	object
51	specz_ra	40914	non-null	float64
52	specz_redshift	40914	non-null	float64
53	specz_redshift_err	40914	non-null	float64
54	x_coord	40914	non-null	float64
55	y_central_image_pop_10px_rad	40914	non-null	int64
56	y_central_image_pop_15px_rad	40914	non-null	int64
57	y_central_image_pop_5px_rad	40914	non-null	int64
58	y_cmodel_mag	40914	non-null	float64
59	y_cmodel_magsigma	40914	non-null	float64
60	y_coord	40914	non-null	float64
61	y_ellipticity	40914	non-null	float64
62	y_half_light_radius	40914	non-null	float64
63	y_isophotal_area	40914	non-null	int64
64	y_major_axis	40914	non-null	float64
65	y_minor_axis	40914	non-null	float64
66	y_peak_surface_brightness	40914	non-null	float64
67	y_petro_rad	40914	non-null	float64
68	y_pos_angle	40914	non-null	float64
69	y_sersic_index	40914	non-null	float64
70	z_central_image_pop_10px_rad	40914	non-null	int64

```

71 z_central_image_pop_15px_rad 40914 non-null int64
72 z_central_image_pop_5px_rad 40914 non-null int64
73 z_cmodel_mag 40914 non-null float64
74 z_cmodel_magsigma 40914 non-null float64
75 z_ellipticity 40914 non-null float64
76 z_half_light_radius 40914 non-null float64
77 z_isophotal_area 40914 non-null int64
78 z_major_axis 40914 non-null float64
79 z_minor_axis 40914 non-null float64
80 z_peak_surface_brightness 40914 non-null float64
81 z_petro_rad 40914 non-null float64
82 z_pos_angle 40914 non-null float64
83 z_sersic_index 40914 non-null float64
84 g_flux 40914 non-null float64
85 r_flux 40914 non-null float64
86 i_flux 40914 non-null float64
87 y_flux 40914 non-null float64
88 z_flux 40914 non-null float64

```

dtypes: bool(1), float64(64), int64(22), object(2)

memory usage: 27.5+ MB

None

	dec	g_central_image_pop_10px_rad \
count	40914.000000	40914.000000
mean	4.772650	1.016278
std	14.789896	0.202830
min	-7.217183	0.000000
25%	-0.960840	1.000000
50%	0.311508	1.000000
75%	1.712535	1.000000
max	53.260936	3.000000

	g_central_image_pop_15px_rad	g_central_image_pop_5px_rad \
count	40914.000000	40914.000000
mean	1.032141	0.998142
std	0.240109	0.160037
min	0.000000	0.000000
25%	1.000000	1.000000
50%	1.000000	1.000000
75%	1.000000	1.000000
max	4.000000	3.000000

	g_cmodel_mag	g_cmodel_magsigma	g_ellipticity	g_half_light_radius
count	40914.000000	40914.000000	40914.000000	40914.000000
mean	21.260192	0.010414	0.230579	6.509069
std	1.882030	0.091981	0.171495	3.625851
min	14.753462	0.000149	0.000000	0.000000
25%	19.733286	0.001531	0.095000	3.905000
50%	21.492703	0.004890	0.190000	5.679000
75%	22.614446	0.011530	0.328000	7.976000
max	30.276308	17.593214	0.896000	26.335000

	g_isophotal_area	g_major_axis	...	z_minor_axis \
count	40914.000000	40914.000000	...	40914.000000
mean	694.781786	5.523996	...	4.064275
std	905.508200	4.095328	...	2.145741
min	0.000000	0.000000	...	0.000000
25%	122.000000	2.706250	...	2.308250
50%	262.000000	3.938000	...	3.740000
75%	947.000000	7.249000	...	5.250000

max	6434.000000	33.320000	...	16.530000	
	z_peak_surface_brightness	z_petro_rad	z_pos_angle	z_sersic_ind	
ex \					
count	40914.000000	40914.000000	40914.000000	40914.000000	40914.000000
mean	-6.819224	6.249483	1.738910	1.6644	
std	1.418684	1.392911	52.237031	0.8703	
min	-10.477100	0.000000	-89.990000	0.0000	
25%	-7.830100	5.280000	-43.900000	1.0370	
50%	-7.000600	5.940000	3.525000	1.5230	
75%	-5.935925	6.600000	46.870000	2.0950	
max	0.000000	10.560000	90.000000	9.9740	

	g_flux	r_flux	i_flux	y_flux	z_flu
x					
count	4.091400e+04	4.091400e+04	4.091400e+04	4.091400e+04	4.091400e+04
mean	1.210340e-08	2.768307e-08	4.473075e-08	6.989669e-08	5.870488e-08
std	2.736938e-08	5.694149e-08	8.773860e-08	1.366205e-07	1.134256e-07
min	7.753125e-13	1.954729e-12	8.507724e-12	3.916208e-11	3.854886e-11
25%	8.999569e-10	2.390220e-09	4.067133e-09	6.001859e-09	5.229988e-09
50%	2.528826e-09	7.571469e-09	1.592592e-08	2.686640e-08	2.237941e-08
75%	1.278456e-08	3.078038e-08	4.926195e-08	7.639837e-08	6.444087e-08
max	1.254918e-06	2.506786e-06	3.731277e-06	5.268088e-06	4.333962e-06

[8 rows x 86 columns]

Clustering (DBScan, K means and Gaussian Mixture Method)

```

In [3]: from sklearn.cluster import DBSCAN, KMeans
        from sklearn.mixture import GaussianMixture
        from sklearn.neighbors import NearestNeighbors
        from sklearn.metrics import silhouette_score
        import matplotlib.pyplot as plt
        import numpy as np
        from kneed import KneeLocator

        # ---- STEP 6: DETERMINE OPTIMAL DBSCAN EPS ----
        k = 5 # Typically, min_samples value
        nearest_neighbors = NearestNeighbors(n_neighbors=k)
        nearest_neighbors.fit(X_umap)
        distances, indices = nearest_neighbors.kneighbors(X_umap)

        # Sort distances to find the "knee" point
        distances = np.sort(distances[:, -1])

```

```

# Use KneeLocator to find optimal epsilon
knee_locator = KneeLocator(range(len(distances)), distances, curve="convex")
optimal_eps = distances[knee_locator.elbow]
print(f"Optimal eps for DBSCAN: {optimal_eps:.3f}")

# ---- STEP 6A: APPLY DBSCAN CLUSTERING WITH OPTIMAL EPS ----
dbscan = DBSCAN(eps=optimal_eps, min_samples=k, metric='euclidean')
cluster_labels_dbscan = dbscan.fit_predict(X_umap)
df_clean.loc[:, 'cluster_dbscan'] = cluster_labels_dbscan

# ---- STEP 6B: APPLY K-MEANS & GMM ----
inertia = []
silhouette_scores = []
k_range = range(2, 10)

for k in k_range:
    kmeans = KMeans(n_clusters=k, random_state=42, n_init='auto')
    labels = kmeans.fit_predict(X_umap)
    inertia.append(kmeans.inertia_)
    score = silhouette_score(X_umap, labels)
    silhouette_scores.append(score)

knee_locator = KneeLocator(k_range, inertia, curve="convex", direction="down")
optimal_k = knee_locator.elbow
print(f"Optimal k for K-Means/GMM: {optimal_k}")

# Apply K-Means with the optimal k
kmeans = KMeans(n_clusters=optimal_k, random_state=42, n_init='auto')
df_clean.loc[:, 'cluster_kmeans'] = kmeans.fit_predict(X_umap)

# Apply Gaussian Mixture Model (GMM)
gmm = GaussianMixture(n_components=optimal_k, random_state=42)
df_clean.loc[:, 'cluster_gmm'] = gmm.fit_predict(X_umap)

# ---- STEP 7: COMPUTE SILHOUETTE SCORES ----
silhouette_kmeans = silhouette_score(X_umap, df_clean['cluster_kmeans'])
silhouette_gmm = silhouette_score(X_umap, df_clean['cluster_gmm'])
silhouette_dbscan = silhouette_score(X_umap, df_clean['cluster_dbscan'])

print(f"Silhouette Scores - KMeans: {silhouette_kmeans:.3f}, GMM: {silhouette_gmm:.3f}, DBSCAN: {silhouette_dbscan:.3f}")

# ---- STEP 8: VISUALIZE CLUSTERING METHODS (2D PLOTS) ----
fig, ax = plt.subplots(1, 3, figsize=(18, 5))

# DBSCAN 2D Plot
ax[0].scatter(X_umap[:, 0], X_umap[:, 1], c=cluster_labels_dbscan, cmap='viridis')
ax[0].set_title("DBSCAN Clustering")
ax[0].set_xlabel("UMAP Component 1")
ax[0].set_ylabel("UMAP Component 2")

# K-Means 2D Plot
ax[1].scatter(X_umap[:, 0], X_umap[:, 1], c=df_clean['cluster_kmeans'], cmap='viridis')
ax[1].set_title("K-Means Clustering")
ax[1].set_xlabel("UMAP Component 1")
ax[1].set_ylabel("UMAP Component 2")

# GMM 2D Plot
ax[2].scatter(X_umap[:, 0], X_umap[:, 1], c=df_clean['cluster_gmm'], cmap='viridis')

```




```
ax[2].set_title("GMM Clustering")
ax[2].set_xlabel("UMAP Component 1")
ax[2].set_ylabel("UMAP Component 2")

plt.show()
```

Optimal eps for DBSCAN: 0.323

Optimal k for K-Means/GMM: 4

Silhouette Scores - KMeans: 0.486, GMM: 0.455, DBSCAN: -0.422



Random Forest

```
In [4]: # ---- STEP 8: RANDOM FOREST REGRESSION ----
# Train separate Random Forest models for each clustering method using fl

results = {}

for cluster_type in ['cluster_dbscan', 'cluster_kmeans', 'cluster_gmm']:
    df_temp = pd.get_dummies(df_clean, columns=[cluster_type], prefix=f'
    X = df_temp[['g_flux', 'r_flux', 'i_flux', 'y_flux', 'z_flux']] +
        [col for col in df_temp.columns if col.startswith(f'clust
    y = df_temp['specz_redshift']

    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0

    rf = RandomForestRegressor(n_estimators=100, random_state=42)
    rf.fit(X_train, y_train)

    y_pred = rf.predict(X_test)

    mse = mean_squared_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)


    results[cluster_type] = {'MSE': mse, 'R2': r2, 'y_pred': y_pred}

# ---- STEP 9: VISUALIZE REGRESSION RESULTS ----
fig, ax = plt.subplots(1, 3, figsize=(18, 6))
titles = ['DBSCAN', 'K-Means', 'GMM']
colors = ['blue', 'green', 'red']

for i, cluster_type in enumerate(results.keys()):
    ax[i].scatter(y_test, results[cluster_type]['y_pred'], alpha=0.6, col
    ax[i].set_xlabel('Actual Redshift')
    ax[i].set_ylabel('Predicted Redshift')
    ax[i].set_title(f'Actual vs Predicted Redshift ({titles[i]})')
    ax[i].plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()])

plt.tight_layout()
plt.show()
```

```
# ---- STEP 10: COMPARE RESULTS ACROSS CLUSTERING METHODS ----
print("Random Forest Regression Results:")
for cluster_type, metrics in results.items():
    print(f"{cluster_type.upper()} - MSE: {metrics['MSE']:.6f}, R2: {metrics['R2']:.6f}")
```



Random Forest Regression Results:

CLUSTER_DBSCAN - MSE: 0.103494, R2: 0.673538

CLUSTER_KMEANS - MSE: 0.103177, R2: 0.674537

CLUSTER_GMM - MSE: 0.103596, R2: 0.673214

Checking

```
In [5]: print(df_clean.head())
print(df_clean.columns)
```

	g_flux	r_flux	i_flux	y_flux	z_flux
0	7.482335e-09	2.987891e-08	5.117533e-08	8.117985e-08	6.737572e-08
1	1.297347e-09	5.643225e-09	1.319162e-08	2.141019e-08	1.828079e-08
2	3.471398e-09	1.571304e-08	3.743521e-08	6.504857e-08	5.585498e-08
3	4.114510e-08	1.090914e-07	1.841788e-07	3.080686e-07	2.402212e-07
4	3.987357e-09	1.373544e-08	2.634882e-08	4.449762e-08	3.527128e-08

	specz_redshift	g_r	r_i	i_y	y_z
0	0.31652	-2.239657e-08	-2.129643e-08	-3.000452e-08	1.380413e-08
1	0.56769	-4.345878e-09	-7.548391e-09	-8.218576e-09	3.129404e-09
2	0.53428	-1.224164e-08	-2.172218e-08	-2.761336e-08	9.193588e-09
3	0.11878	-6.794635e-08	-7.508739e-08	-1.238898e-07	6.784739e-08
4	0.23497	-9.748082e-09	-1.261338e-08	-1.814880e-08	9.226337e-09

	Cluster	cluster_dbscan	cluster_kmeans	cluster_gmm
0	2	0	3	3
1	1	0	1	1
2	1	0	3	3
3	2	0	2	2
4	1	0	3	3

```
Index(['g_flux', 'r_flux', 'i_flux', 'y_flux', 'z_flux', 'specz_redshift',
      'g_r', 'r_i', 'i_y', 'y_z', 'Cluster', 'cluster_dbscan',
      'cluster_kmeans', 'cluster_gmm'],
      dtype='object')
```

SVR and UMAP

```
In [6]: # Add-Ons for Photometric Redshift Estimation
# Fixing UMAP Import Issue and Enhancing Preprocessing

# ---- STEP 11: Install and Import UMAP Properly ----
```

```

# Ensure proper UMAP installation: pip install umap-learn
from umap import UMAP

umap = UMAP(n_neighbors=15, min_dist=0.1, n_components=2, random_state=42)
X_umap = umap.fit_transform(df_transformed)

# ---- STEP 12: Add More Color Indices (Using Flux Values) ----
df_clean['u_g'] = df_clean['g_flux'] - df_clean['r_flux']
df_clean['g_r'] = df_clean['r_flux'] - df_clean['i_flux']
df_clean['r_i'] = df_clean['i_flux'] - df_clean['y_flux']
df_clean['i_z'] = df_clean['y_flux'] - df_clean['z_flux']

# ---- STEP 13: Additional Clustering (HDBSCAN) for SVR Comparison ----
try:
    from hdbscan import HDBSCAN
except ImportError:
    print("HDBSCAN not installed. Use: pip install hdbscan")
    raise

hdbscan = HDBSCAN(min_cluster_size=8)
df_clean['cluster_hdbscan'] = hdbscan.fit_predict(X_umap)

# ---- STEP 14: SVR Regression for Redshift Estimation ----
from sklearn.svm import SVR
svr_results = {}

for cluster_type in ['cluster_kmeans', 'cluster_gmm', 'cluster_hdbscan']:
    df_temp = pd.get_dummies(df_clean, columns=[cluster_type], prefix=[f'
    X = df_temp.drop(columns=['specz_redshift'])
    y = df_temp['specz_redshift']

    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0

    svr = SVR(kernel='rbf', C=1.0, epsilon=0.1)
    svr.fit(X_train, y_train)
    y_pred = svr.predict(X_test)

    mse = mean_squared_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)

    svr_results[cluster_type] = {'MSE': mse, 'R2': r2, 'y_pred': y_pred}

# ---- STEP 15: Plot SVR Results ----
fig, ax = plt.subplots(1, 3, figsize=(18, 6))
titles = ['K-Means', 'GMM', 'HDBSCAN']
colors = ['purple', 'orange', 'cyan']

for i, cluster_type in enumerate(svr_results.keys()):
    ax[i].scatter(y_test, svr_results[cluster_type]['y_pred'], alpha=0.6,
    ax[i].set_title(f'SVR Actual vs Predicted ({titles[i]}')
    ax[i].set_xlabel('Actual Redshift')
    ax[i].set_ylabel('Predicted Redshift')
    ax[i].plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()])

plt.tight_layout()
plt.show()

# ---- Print SVR Performance Metrics ----
print('SVR Regression Results:')

```

```
for cluster_type, metrics in svr_results.items():
    print(f"{cluster_type.upper()} - MSE: {metrics['MSE']:.6f}, R2: {metrics['R2']:.6f}")
```

```
/home/chloy/miniconda3/lib/python3.10/site-packages/sklearn/utils/deprecation.py:151: FutureWarning: 'force_all_finite' was renamed to 'ensure_all_finite' in 1.6 and will be removed in 1.8.
```

```
warnings.warn(
```

```
/home/chloy/miniconda3/lib/python3.10/site-packages/umap/umap_.py:1952: UserWarning: n_jobs value 1 overridden to 1 by setting random_state. Use no seed for parallelism.
```


```
warn(
```

```
/home/chloy/miniconda3/lib/python3.10/site-packages/sklearn/utils/deprecation.py:151: FutureWarning: 'force_all_finite' was renamed to 'ensure_all_finite' in 1.6 and will be removed in 1.8.
```

```
warnings.warn(
```

```
/home/chloy/miniconda3/lib/python3.10/site-packages/sklearn/utils/deprecation.py:151: FutureWarning: 'force_all_finite' was renamed to 'ensure_all_finite' in 1.6 and will be removed in 1.8.
```

```
warnings.warn(
```



SVR Regression Results:

CLUSTER_KMEANS - MSE: 0.227899, R2: 0.281111

CLUSTER_GMM - MSE: 0.226127, R2: 0.286701

CLUSTER_HDBSCAN - MSE: 0.167378, R2: 0.472019