

Hybrid_Model

February 24, 2025

Preparing Data And Preprocessing

```
[2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.ensemble import RandomForestRegressor
from sklearn.decomposition import PCA
from sklearn.cluster import DBSCAN
from sklearn.preprocessing import StandardScaler, PowerTransformer
from sklearn.feature_selection import VarianceThreshold
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.neighbors import LocalOutlierFactor
from kneed import KneeLocator
from sklearn.cluster import KMeans
from sklearn.mixture import GaussianMixture
import umap.umap_ as umap

# ---- STEP 1: LOAD DATA ----
df = pd.read_excel('./data_files/Data_re.xlsx') # Ensure the file exists

# ---- STEP 2: FEATURE SELECTION ----
cols_to_keep = ['g_flux', 'r_flux', 'i_flux', 'y_flux', 'z_flux',
               'specz_redshift']
df_selected = df[cols_to_keep].copy()

# Compute Flux Color Indices
df_selected['g_r'] = df_selected['g_flux'] - df_selected['r_flux']
df_selected['r_i'] = df_selected['r_flux'] - df_selected['i_flux']
df_selected['i_y'] = df_selected['i_flux'] - df_selected['y_flux']
df_selected['y_z'] = df_selected['y_flux'] - df_selected['z_flux']

# ---- STEP 3: REMOVE OUTLIERS ----
lof = LocalOutlierFactor(n_neighbors=20, contamination=0.02)
outlier_scores = lof.fit_predict(df_selected.drop(columns=['specz_redshift']))
df_clean = df_selected[outlier_scores == 1].copy()
```

```

# ---- STEP 4: SCALING & TRANSFORMING ----
scaler = StandardScaler()
df_scaled = scaler.fit_transform(df_clean.drop(columns=['specz_redshift']))

# Apply Power Transformation (Yeo-Johnson for normalizing skewed data)
power_transformer = PowerTransformer(method='yeo-johnson')
df_transformed = power_transformer.fit_transform(df_scaled)

# ---- STEP 5: DIMENSIONALITY REDUCTION WITH UMAP ----
reducer = umap.UMAP(n_components=2, n_neighbors=15, min_dist=0.1,
↳random_state=42)
X_umap = reducer.fit_transform(df_transformed)

# ---- STEP 6: OPTIONAL PCA FOR EXPLORATION ----
pca = PCA(n_components=3)
X_pca = pca.fit_transform(df_transformed)

# ---- STEP 7: CLUSTERING ----
kmeans = KMeans(n_clusters=3, random_state=42)
df_clean['Cluster'] = kmeans.fit_predict(X_umap)

# ---- PLOT RESULTS ----
plt.scatter(X_umap[:, 0], X_umap[:, 1], c=df_clean['Cluster'], cmap='viridis')
plt.xlabel('UMAP Component 1')
plt.ylabel('UMAP Component 2')
plt.title('UMAP Clustering')
plt.show()


```

```

2025-02-24 21:14:13.144587: E
external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:477] Unable to register
cuFFT factory: Attempting to register factory for plugin cuFFT when one has
already been registered
WARNING: All log messages before absl::InitializeLog() is called are written to
STDERR
E0000 00:00:1740411853.199942    3391 cuda_dnn.cc:8310] Unable to register cuDNN
factory: Attempting to register factory for plugin cuDNN when one has already
been registered
E0000 00:00:1740411853.215569    3391 cuda_blas.cc:1418] Unable to register
cuBLAS factory: Attempting to register factory for plugin cuBLAS when one has
already been registered
2025-02-24 21:14:13.331174: I tensorflow/core/platform/cpu_feature_guard.cc:210]
This TensorFlow binary is optimized to use available CPU instructions in
performance-critical operations.
To enable the following instructions: AVX2 FMA, in other operations, rebuild
TensorFlow with the appropriate compiler flags.
/home/chloy/miniconda3/lib/python3.10/site-
packages/sklearn/utils/deprecation.py:151: FutureWarning: 'force_all_finite' was
renamed to 'ensure_all_finite' in 1.6 and will be removed in 1.8.

```

```
warnings.warn(
/home/chloy/miniconda3/lib/python3.10/site-packages/umap/umap_.py:1952:
UserWarning: n_jobs value 1 overridden to 1 by setting random_state. Use no seed
for parallelism.
warn(
```



checking

```
[3]: print(df.head())
      print(df.info())
      print(df.describe())
```

```

                                coord      dec  \
0  b'(179325.3125, 99694.8046875, -21178.96484375)' -5.893433
1  b'(179236.609375, 99349.8203125, -23431.181640... -6.522742
2    b'(179281.5, 99283.5078125, -23368.759765625)' -6.505290
3  b'(179365.171875, 99158.046875, -23259.0839843... -6.474628
4    b'(179366.421875, 99172.25, -23188.84765625)' -6.454993

g_central_image_pop_10px_rad  g_central_image_pop_15px_rad  \
0                               1                               1
```

1	1	1
2	1	1
3	1	1
4	1	1

	g_central_image_pop_5px_rad	g_cmodel_mag	g_cmodel_magsigma \
0	1	20.314907	0.002624
1	1	22.217360	0.010902
2	1	21.148739	0.008013
3	1	18.464205	0.001740
4	1	20.998287	0.006011

	g_ellipticity	g_half_light_radius	g_isophotal_area ...	z_minor_axis \
0	0.147	6.047	603 ...	4.938
1	0.130	3.430	93 ...	2.713
2	0.209	6.597	254 ...	4.351
3	0.525	10.855	1064 ...	4.815
4	0.738	8.261	386 ...	2.279

	z_peak_surface_brightness	z_petro_rad	z_pos_angle	z_sersic_index \
0	-8.2933	5.28	36.15	2.193
1	-7.3657	5.94	-61.78	1.649
2	-7.6539	9.24	32.76	2.364
3	-8.5825	6.60	53.15	1.494
4	-6.8798	5.94	21.16	1.063

	g_flux	r_flux	i_flux	y_flux	z_flux
0	7.482335e-09	2.987891e-08	5.117533e-08	8.117985e-08	6.737572e-08
1	1.297347e-09	5.643225e-09	1.319162e-08	2.141019e-08	1.828079e-08
2	3.471398e-09	1.571304e-08	3.743521e-08	6.504857e-08	5.585498e-08
3	4.114510e-08	1.090914e-07	1.841788e-07	3.080686e-07	2.402212e-07
4	3.987357e-09	1.373544e-08	2.634882e-08	4.449762e-08	3.527128e-08

[5 rows x 89 columns]

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 40914 entries, 0 to 40913

Data columns (total 89 columns):

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	coord	40914 non-null	object
1	dec	40914 non-null	float64
2	g_central_image_pop_10px_rad	40914 non-null	int64
3	g_central_image_pop_15px_rad	40914 non-null	int64
4	g_central_image_pop_5px_rad	40914 non-null	int64
5	g_cmodel_mag	40914 non-null	float64
6	g_cmodel_magsigma	40914 non-null	float64
7	g_ellipticity	40914 non-null	float64
8	g_half_light_radius	40914 non-null	float64

9	g_isophotal_area	40914	non-null	int64
10	g_major_axis	40914	non-null	float64
11	g_minor_axis	40914	non-null	float64
12	g_peak_surface_brightness	40914	non-null	float64
13	g_petro_rad	40914	non-null	float64
14	g_pos_angle	40914	non-null	float64
15	g_sersic_index	40914	non-null	float64
16	i_central_image_pop_10px_rad	40914	non-null	int64
17	i_central_image_pop_15px_rad	40914	non-null	int64
18	i_central_image_pop_5px_rad	40914	non-null	int64
19	i_cmodel_mag	40914	non-null	float64
20	i_cmodel_magsigma	40914	non-null	float64
21	i_ellipticity	40914	non-null	float64
22	i_half_light_radius	40914	non-null	float64
23	i_isophotal_area	40914	non-null	int64
24	i_major_axis	40914	non-null	float64
25	i_minor_axis	40914	non-null	float64
26	i_peak_surface_brightness	40914	non-null	float64
27	i_petro_rad	40914	non-null	float64
28	i_pos_angle	40914	non-null	float64
29	i_sersic_index	40914	non-null	float64
30	object_id	40914	non-null	int64
31	r_central_image_pop_10px_rad	40914	non-null	int64
32	r_central_image_pop_15px_rad	40914	non-null	int64
33	r_central_image_pop_5px_rad	40914	non-null	int64
34	r_cmodel_mag	40914	non-null	float64
35	r_cmodel_magsigma	40914	non-null	float64
36	r_ellipticity	40914	non-null	float64
37	r_half_light_radius	40914	non-null	float64
38	r_isophotal_area	40914	non-null	int64
39	r_major_axis	40914	non-null	float64
40	r_minor_axis	40914	non-null	float64
41	r_peak_surface_brightness	40914	non-null	float64
42	r_petro_rad	40914	non-null	float64
43	r_pos_angle	40914	non-null	float64
44	r_sersic_index	40914	non-null	float64
45	ra	40914	non-null	float64
46	skymap_id	40914	non-null	int64
47	specz_dec	40914	non-null	float64
48	specz_flag_homogeneous	40914	non-null	bool
49	specz_mag_i	40914	non-null	float64
50	specz_name	40914	non-null	object
51	specz_ra	40914	non-null	float64
52	specz_redshift	40914	non-null	float64
53	specz_redshift_err	40914	non-null	float64
54	x_coord	40914	non-null	float64
55	y_central_image_pop_10px_rad	40914	non-null	int64
56	y_central_image_pop_15px_rad	40914	non-null	int64

57	y_central_image_pop_5px_rad	40914	non-null	int64
58	y_cmodel_mag	40914	non-null	float64
59	y_cmodel_magsigma	40914	non-null	float64
60	y_coord	40914	non-null	float64
61	y_ellipticity	40914	non-null	float64
62	y_half_light_radius	40914	non-null	float64
63	y_isophotal_area	40914	non-null	int64
64	y_major_axis	40914	non-null	float64
65	y_minor_axis	40914	non-null	float64
66	y_peak_surface_brightness	40914	non-null	float64
67	y_petro_rad	40914	non-null	float64
68	y_pos_angle	40914	non-null	float64
69	y_sersic_index	40914	non-null	float64
70	z_central_image_pop_10px_rad	40914	non-null	int64
71	z_central_image_pop_15px_rad	40914	non-null	int64
72	z_central_image_pop_5px_rad	40914	non-null	int64
73	z_cmodel_mag	40914	non-null	float64
74	z_cmodel_magsigma	40914	non-null	float64
75	z_ellipticity	40914	non-null	float64
76	z_half_light_radius	40914	non-null	float64
77	z_isophotal_area	40914	non-null	int64
78	z_major_axis	40914	non-null	float64
79	z_minor_axis	40914	non-null	float64
80	z_peak_surface_brightness	40914	non-null	float64
81	z_petro_rad	40914	non-null	float64
82	z_pos_angle	40914	non-null	float64
83	z_sersic_index	40914	non-null	float64
84	g_flux	40914	non-null	float64
85	r_flux	40914	non-null	float64
86	i_flux	40914	non-null	float64
87	y_flux	40914	non-null	float64
88	z_flux	40914	non-null	float64

dtypes: bool(1), float64(64), int64(22), object(2)

memory usage: 27.5+ MB

None

	dec	g_central_image_pop_10px_rad	\
count	40914.000000	40914.000000	
mean	4.772650	1.016278	
std	14.789896	0.202830	
min	-7.217183	0.000000	
25%	-0.960840	1.000000	
50%	0.311508	1.000000	
75%	1.712535	1.000000	
max	53.260936	3.000000	

	g_central_image_pop_15px_rad	g_central_image_pop_5px_rad	\
count	40914.000000	40914.000000	
mean	1.032141	0.998142	

std	0.240109	0.160037
min	0.000000	0.000000
25%	1.000000	1.000000
50%	1.000000	1.000000
75%	1.000000	1.000000
max	4.000000	3.000000

	g_cmodel_mag	g_cmodel_magsigma	g_ellipticity	g_half_light_radius \
count	40914.000000	40914.000000	40914.000000	40914.000000
mean	21.260192	0.010414	0.230579	6.509069
std	1.882030	0.091981	0.171495	3.625851
min	14.753462	0.000149	0.000000	0.000000
25%	19.733286	0.001531	0.095000	3.905000
50%	21.492703	0.004890	0.190000	5.679000
75%	22.614446	0.011530	0.328000	7.976000
max	30.276308	17.593214	0.896000	26.335000

	g_isophotal_area	g_major_axis	...	z_minor_axis \
count	40914.000000	40914.000000	...	40914.000000
mean	694.781786	5.523996	...	4.064275
std	905.508200	4.095328	...	2.145741
min	0.000000	0.000000	...	0.000000
25%	122.000000	2.706250	...	2.308250
50%	262.000000	3.938000	...	3.740000
75%	947.000000	7.249000	...	5.250000
max	6434.000000	33.320000	...	16.530000

	z_peak_surface_brightness	z_petro_rad	z_pos_angle	z_sersic_index \
count	40914.000000	40914.000000	40914.000000	40914.000000
mean	-6.819224	6.249483	1.738910	1.664449
std	1.418684	1.392911	52.237031	0.870328
min	-10.477100	0.000000	-89.990000	0.000000
25%	-7.830100	5.280000	-43.900000	1.037000
50%	-7.000600	5.940000	3.525000	1.523000
75%	-5.935925	6.600000	46.870000	2.095000
max	0.000000	10.560000	90.000000	9.974000

	g_flux	r_flux	i_flux	y_flux	z_flux
count	4.091400e+04	4.091400e+04	4.091400e+04	4.091400e+04	4.091400e+04
mean	1.210340e-08	2.768307e-08	4.473075e-08	6.989669e-08	5.870488e-08
std	2.736938e-08	5.694149e-08	8.773860e-08	1.366205e-07	1.134256e-07
min	7.753125e-13	1.954729e-12	8.507724e-12	3.916208e-11	3.854886e-11
25%	8.999569e-10	2.390220e-09	4.067133e-09	6.001859e-09	5.229988e-09
50%	2.528826e-09	7.571469e-09	1.592592e-08	2.686640e-08	2.237941e-08
75%	1.278456e-08	3.078038e-08	4.926195e-08	7.639837e-08	6.444087e-08
max	1.254918e-06	2.506786e-06	3.731277e-06	5.268088e-06	4.333962e-06

[8 rows x 86 columns]

Clustering (DBSCAN, K means and Gaussian Mixture Method)

```
[4]: from sklearn.cluster import DBSCAN, KMeans
from sklearn.mixture import GaussianMixture
from sklearn.neighbors import NearestNeighbors
from sklearn.metrics import silhouette_score
import matplotlib.pyplot as plt
import numpy as np
from kneed import KneeLocator

# ---- STEP 6: DETERMINE OPTIMAL DBSCAN EPS ----
k = 5 # Typically, min_samples value
nearest_neighbors = NearestNeighbors(n_neighbors=k)
nearest_neighbors.fit(X_umap)
distances, indices = nearest_neighbors.kneighbors(X_umap)

# Sort distances to find the "knee" point
distances = np.sort(distances[:, -1])

# Use KneeLocator to find optimal epsilon
knee_locator = KneeLocator(range(len(distances)), distances, curve="convex",
    ↪direction="increasing")
optimal_eps = distances[knee_locator.elbow]
print(f"Optimal eps for DBSCAN: {optimal_eps:.3f}")

# ---- STEP 6A: APPLY DBSCAN CLUSTERING WITH OPTIMAL EPS ----
dbscan = DBSCAN(eps=optimal_eps, min_samples=k, metric='euclidean')
cluster_labels_dbscan = dbscan.fit_predict(X_umap)
df_clean.loc[:, 'cluster_dbscan'] = cluster_labels_dbscan

# ---- STEP 6B: APPLY K-MEANS & GMM ----
inertia = []
silhouette_scores = []
k_range = range(2, 10)

for k in k_range:
    kmeans = KMeans(n_clusters=k, random_state=42, n_init='auto')
    labels = kmeans.fit_predict(X_umap)
    inertia.append(kmeans.inertia_)
    score = silhouette_score(X_umap, labels)
    silhouette_scores.append(score)

knee_locator = KneeLocator(k_range, inertia, curve="convex",
    ↪direction="decreasing")
optimal_k = knee_locator.elbow
print(f"Optimal k for K-Means/GMM: {optimal_k}")
```



```

# Apply K-Means with the optimal k
kmeans = KMeans(n_clusters=optimal_k, random_state=42, n_init='auto')
df_clean.loc[:, 'cluster_kmeans'] = kmeans.fit_predict(X_umap)

# Apply Gaussian Mixture Model (GMM)
gmm = GaussianMixture(n_components=optimal_k, random_state=42)
df_clean.loc[:, 'cluster_gmm'] = gmm.fit_predict(X_umap)

# ---- STEP 7: COMPUTE SILHOUETTE SCORES ----
silhouette_kmeans = silhouette_score(X_umap, df_clean['cluster_kmeans'])
silhouette_gmm = silhouette_score(X_umap, df_clean['cluster_gmm'])
silhouette_dbscan = silhouette_score(X_umap[df_clean['cluster_dbscan'] != -1],
                                     df_clean['cluster_dbscan'][df_clean['cluster_dbscan'] != -1])

print(f"Silhouette Scores - KMeans: {silhouette_kmeans:.3f}, GMM: {silhouette_gmm:.3f}, DBSCAN: {silhouette_dbscan:.3f}")

# ---- STEP 8: VISUALIZE CLUSTERING METHODS (2D PLOTS) ----
fig, ax = plt.subplots(1, 3, figsize=(18, 5))

# DBSCAN 2D Plot
ax[0].scatter(X_umap[:, 0], X_umap[:, 1], c=cluster_labels_dbscan,
              cmap='viridis', alpha=0.6)
ax[0].set_title("DBSCAN Clustering")
ax[0].set_xlabel("UMAP Component 1")
ax[0].set_ylabel("UMAP Component 2")

# K-Means 2D Plot
ax[1].scatter(X_umap[:, 0], X_umap[:, 1], c=df_clean['cluster_kmeans'],
              cmap='viridis', alpha=0.6)
ax[1].set_title("K-Means Clustering")
ax[1].set_xlabel("UMAP Component 1")
ax[1].set_ylabel("UMAP Component 2")

# GMM 2D Plot
ax[2].scatter(X_umap[:, 0], X_umap[:, 1], c=df_clean['cluster_gmm'],
              cmap='viridis', alpha=0.6)
ax[2].set_title("GMM Clustering")
ax[2].set_xlabel("UMAP Component 1")
ax[2].set_ylabel("UMAP Component 2")


plt.show()

```

Optimal eps for DBSCAN: 0.323

Optimal k for K-Means/GMM: 4

Silhouette Scores - KMeans: 0.486, GMM: 0.455, DBSCAN: -0.422



Random Forest

```
[5]: # ---- STEP 8: RANDOM FOREST REGRESSION ----
# Train separate Random Forest models for each clustering method using flux_
# features

results = {}

for cluster_type in ['cluster_dbscan', 'cluster_kmeans', 'cluster_gmm']:
    df_temp = pd.get_dummies(df_clean, columns=[cluster_type],
    prefix=[f'clust_{cluster_type}'])
    X = df_temp[['g_flux', 'r_flux', 'i_flux', 'y_flux', 'z_flux'] +
    [col for col in df_temp.columns if col.
    startswith(f'clust_{cluster_type}')]]
    y = df_temp['specz_redshift']

    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

    rf = RandomForestRegressor(n_estimators=100, random_state=42)
    rf.fit(X_train, y_train)

    y_pred = rf.predict(X_test)

    mse = mean_squared_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)

    results[cluster_type] = {'MSE': mse, 'R2': r2, 'y_pred': y_pred}

# ---- STEP 9: VISUALIZE REGRESSION RESULTS ----
fig, ax = plt.subplots(1, 3, figsize=(18, 6))
titles = ['DBSCAN', 'K-Means', 'GMM']
colors = ['blue', 'green', 'red']
```


```

for i, cluster_type in enumerate(results.keys()):
    ax[i].scatter(y_test, results[cluster_type]['y_pred'], alpha=0.6,
        color=colors[i])
    ax[i].set_xlabel('Actual Redshift')
    ax[i].set_ylabel('Predicted Redshift')
    ax[i].set_title(f'Actual vs Predicted Redshift ({titles[i]})')
    ax[i].plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()],
        'k--', lw=2)

plt.tight_layout()
plt.show()

# ---- STEP 10: COMPARE RESULTS ACROSS CLUSTERING METHODS ----
print("Random Forest Regression Results:")
for cluster_type, metrics in results.items():
    print(f"{cluster_type.upper()} - MSE: {metrics['MSE']:.6f}, R2: {metrics['R2']:.6f}")

```



Random Forest Regression Results:

CLUSTER_DBSCAN - MSE: 0.103494, R2: 0.673538

CLUSTER_KMEANS - MSE: 0.103177, R2: 0.674537

CLUSTER_GMM - MSE: 0.103596, R2: 0.673214

Checking

```

[6]: print(df_clean.head())
     print(df_clean.columns)

```

	g_flux	r_flux	i_flux	y_flux	z_flux	\
0	7.482335e-09	2.987891e-08	5.117533e-08	8.117985e-08	6.737572e-08	
1	1.297347e-09	5.643225e-09	1.319162e-08	2.141019e-08	1.828079e-08	
2	3.471398e-09	1.571304e-08	3.743521e-08	6.504857e-08	5.585498e-08	
3	4.114510e-08	1.090914e-07	1.841788e-07	3.080686e-07	2.402212e-07	
4	3.987357e-09	1.373544e-08	2.634882e-08	4.449762e-08	3.527128e-08	

	specz_redshift	g_r	r_i	i_y	y_z \
0	0.31652	-2.239657e-08	-2.129643e-08	-3.000452e-08	1.380413e-08
1	0.56769	-4.345878e-09	-7.548391e-09	-8.218576e-09	3.129404e-09
2	0.53428	-1.224164e-08	-2.172218e-08	-2.761336e-08	9.193588e-09
3	0.11878	-6.794635e-08	-7.508739e-08	-1.238898e-07	6.784739e-08
4	0.23497	-9.748082e-09	-1.261338e-08	-1.814880e-08	9.226337e-09

	Cluster	cluster_dbscan	cluster_kmeans	cluster_gmm
0	2	0	3	3
1	1	0	1	1
2	1	0	3	3
3	2	0	2	2
4	1	0	3	3

```
Index(['g_flux', 'r_flux', 'i_flux', 'y_flux', 'z_flux', 'specz_redshift',
      'g_r', 'r_i', 'i_y', 'y_z', 'Cluster', 'cluster_dbscan',
      'cluster_kmeans', 'cluster_gmm'],
      dtype='object')
```

SVR and UMAP

```
[7]: # Add-Ons for Photometric Redshift Estimation
# Fixing UMAP Import Issue and Enhancing Preprocessing

# ---- STEP 11: Install and Import UMAP Properly ----
# Ensure proper UMAP installation: pip install umap-learn
from umap import UMAP

umap = UMAP(n_neighbors=15, min_dist=0.1, n_components=2, random_state=42)
X_umap = umap.fit_transform(df_transformed)

# ---- STEP 12: Add More Color Indices (Using Flux Values) ----
df_clean['u_g'] = df_clean['g_flux'] - df_clean['r_flux']
df_clean['g_r'] = df_clean['r_flux'] - df_clean['i_flux']
df_clean['r_i'] = df_clean['i_flux'] - df_clean['y_flux']
df_clean['i_z'] = df_clean['y_flux'] - df_clean['z_flux']

# ---- STEP 13: Additional Clustering (HDBSCAN) for SVR Comparison ----
try:
    from hdbscan import HDBSCAN
except ImportError:
    print("HDBSCAN not installed. Use: pip install hdbscan")
    raise

hdbscan = HDBSCAN(min_cluster_size=8)
df_clean['cluster_hdbscan'] = hdbscan.fit_predict(X_umap)

# ---- STEP 14: SVR Regression for Redshift Estimation ----
from sklearn.svm import SVR
```

```

svr_results = {}

for cluster_type in ['cluster_kmeans', 'cluster_gmm', 'cluster_hdbscan']:
    df_temp = pd.get_dummies(df_clean, columns=[cluster_type],
    prefix=[f'clust_{cluster_type}'])
    X = df_temp.drop(columns=['specz_redshift'])
    y = df_temp['specz_redshift']

    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

    svr = SVR(kernel='rbf', C=1.0, epsilon=0.1)
    svr.fit(X_train, y_train)
    y_pred = svr.predict(X_test)

    mse = mean_squared_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)

    svr_results[cluster_type] = {'MSE': mse, 'R2': r2, 'y_pred': y_pred}

# ---- STEP 15: Plot SVR Results ----
fig, ax = plt.subplots(1, 3, figsize=(18, 6))
titles = ['K-Means', 'GMM', 'HDBSCAN']
colors = ['purple', 'orange', 'cyan']

for i, cluster_type in enumerate(svr_results.keys()):
    ax[i].scatter(y_test, svr_results[cluster_type]['y_pred'], alpha=0.6,
    color=colors[i])
    ax[i].set_title(f'SVR Actual vs Predicted ({titles[i]}')
    ax[i].set_xlabel('Actual Redshift')
    ax[i].set_ylabel('Predicted Redshift')
    ax[i].plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()],
    'k--', lw=2)

plt.tight_layout()
plt.show()

# ---- Print SVR Performance Metrics ----
print('SVR Regression Results:')
for cluster_type, metrics in svr_results.items():
    print(f"{cluster_type.upper()} - MSE: {metrics['MSE']:.6f}, R2:
    {metrics['R2']:.6f}")

```

```

/home/chloy/miniconda3/lib/python3.10/site-
packages/sklearn/utils/deprecation.py:151: FutureWarning: 'force_all_finite' was
renamed to 'ensure_all_finite' in 1.6 and will be removed in 1.8.
warnings.warn(


```

```
/home/chloy/miniconda3/lib/python3.10/site-packages/umap/umap_.py:1952:
UserWarning: n_jobs value 1 overridden to 1 by setting random_state. Use no seed
for parallelism.
```

```
warn(
/home/chloy/miniconda3/lib/python3.10/site-
packages/sklearn/utils/deprecation.py:151: FutureWarning: 'force_all_finite' was
renamed to 'ensure_all_finite' in 1.6 and will be removed in 1.8.
```

```
warnings.warn(
/home/chloy/miniconda3/lib/python3.10/site-
packages/sklearn/utils/deprecation.py:151: FutureWarning: 'force_all_finite' was
renamed to 'ensure_all_finite' in 1.6 and will be removed in 1.8.
```

```
warnings.warn(
```



SVR Regression Results:

CLUSTER_KMEANS - MSE: 0.227899, R2: 0.281111

CLUSTER_GMM - MSE: 0.226127, R2: 0.286701

CLUSTER_HDBSCAN - MSE: 0.167378, R2: 0.472019

XGBoost , Gradient boosting and MLP NN

```
[8]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.ensemble import GradientBoostingRegressor
from xgboost import XGBRegressor
from sklearn.neural_network import MLPRegressor
import faiss # Faster nearest neighbor search

# ---- STEP 1: ADD MORE COLOR INDICES ----
# Verify flux columns exist before calculating new features
if all(col in df_clean.columns for col in ['g_flux', 'r_flux', 'i_flux',
↳ 'y_flux', 'z_flux']):
    df_clean['g_i'] = df_clean['g_flux'] - df_clean['i_flux']
```

```

    df_clean['r_z'] = df_clean['r_flux'] - df_clean['z_flux']
    df_clean['i_y_z'] = df_clean['i_flux'] - df_clean['y_flux'] -
    df_clean['z_flux']
else:
    raise ValueError("Missing required flux columns in df_clean.")

# ---- STEP 2: DEFINE FEATURES & TARGET ----
features = ['g_flux', 'r_flux', 'i_flux', 'y_flux', 'z_flux',
            'g_r', 'r_i', 'i_y', 'y_z', 'g_i', 'r_z', 'i_y_z']

# Include cluster one-hot encoding for different clustering techniques
if 'cluster_kmeans' in df_clean.columns and 'cluster_gmm' in df_clean.columns
    and 'cluster_hdbscan' in df_clean.columns:
    df_encoded = pd.get_dummies(df_clean, columns=['cluster_kmeans',
    'cluster_gmm', 'cluster_hdbscan'],
                                prefix=['clust_kmeans', 'clust_gmm',
    'clust_hdbscan'])
else:
    raise ValueError("Cluster columns are missing in df_clean.")

X = df_encoded[features + [col for col in df_encoded.columns if col.
    startswith('clust_')]]
y = df_encoded['specz_redshift']

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

# ---- STEP 3: APPLY XGBOOST ----
xgb = XGBRegressor(n_estimators=200, learning_rate=0.1, random_state=42)
xgb.fit(X_train, y_train)
y_pred_xgb = xgb.predict(X_test)

mse_xgb = mean_squared_error(y_test, y_pred_xgb)
r2_xgb = r2_score(y_test, y_pred_xgb)

print(f"XGBoost - MSE: {mse_xgb:.6f}, R2: {r2_xgb:.6f}")

# ---- STEP 4: APPLY GRADIENT BOOSTING ----
gbr = GradientBoostingRegressor(n_estimators=150, learning_rate=0.05,
    random_state=42)
gbr.fit(X_train, y_train)
y_pred_gbr = gbr.predict(X_test)

mse_gbr = mean_squared_error(y_test, y_pred_gbr)
r2_gbr = r2_score(y_test, y_pred_gbr)

```

```

print(f"Gradient Boosting - MSE: {mse_gbr:.6f}, R2: {r2_gbr:.6f}")

# ---- STEP 5: TRY MLP NEURAL NETWORK ----
mlp = MLPRegressor(hidden_layer_sizes=(64, 32), max_iter=500, random_state=42)
mlp.fit(X_train, y_train)
y_pred_mlp = mlp.predict(X_test)

mse_mlp = mean_squared_error(y_test, y_pred_mlp)
r2_mlp = r2_score(y_test, y_pred_mlp)

print(f"MLP Neural Net - MSE: {mse_mlp:.6f}, R2: {r2_mlp:.6f}")

# ---- STEP 6: OPTIONAL - FASTER NEAREST NEIGHBOR SEARCH FOR DBSCAN ----
index = faiss.IndexFlatL2(X_train.shape[1]) # L2 distance (Euclidean)
index.add(X_train.astype('float32')) # FAISS requires float32
_, indices = index.search(X_train.astype('float32'), k=5) # Find 5 nearest
↳ neighbors

```

XGBoost - MSE: 0.099036, R2: 0.687598
 Gradient Boosting - MSE: 0.112597, R2: 0.644823
 MLP Neural Net - MSE: 0.154431, R2: 0.512862

```

[9]: # ---- PLOT COMPARISON ----
models = ['XGBoost', 'Gradient Boosting', 'MLP Neural Net']
mse_scores = [mse_xgb, mse_gbr, mse_mlp]
r2_scores = [r2_xgb, r2_gbr, r2_mlp]

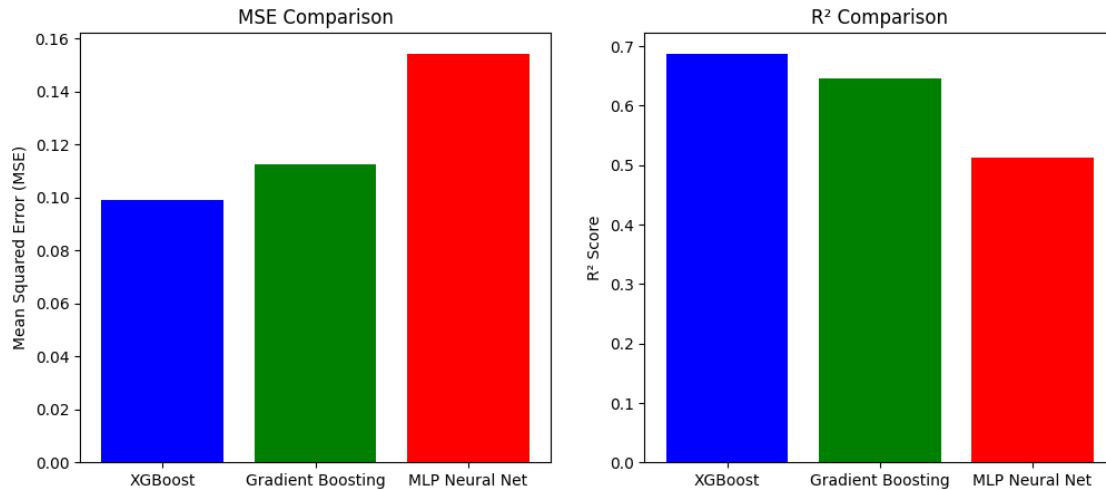
fig, ax = plt.subplots(1, 2, figsize=(12, 5))

# MSE Plot
ax[0].bar(models, mse_scores, color=['blue', 'green', 'red'])
ax[0].set_ylabel("Mean Squared Error (MSE)")
ax[0].set_title("MSE Comparison")

# R2 Score Plot
ax[1].bar(models, r2_scores, color=['blue', 'green', 'red'])
ax[1].set_ylabel("R2 Score")
ax[1].set_title("R2 Comparison")

plt.show()

```

Fine Tuning XGBoost Hyper-Parameters

```
[10]: from xgboost import XGBRegressor
from sklearn.model_selection import GridSearchCV

# ---- XGBOOST HYPERPARAMETER TUNING ----
xgb_params = {
    'n_estimators': [200, 300, 500], # More trees + better accuracy, but slower
    'learning_rate': [0.01, 0.05, 0.1], # Lower = slower but more stable
    'max_depth': [5, 7, 9], # Deeper trees can learn more but risk overfitting
    'subsample': [0.8, 1.0], # Prevents overfitting by using a fraction of
    ↪ data per tree
    'colsample_bytree': [0.8, 1.0] # Randomly selects features per tree to
    ↪ prevent overfitting
}

xgb = XGBRegressor(random_state=42)
xgb_grid = GridSearchCV(xgb, xgb_params, cv=3, scoring='r2', verbose=2,
    ↪ n_jobs=-1)
xgb_grid.fit(X_train, y_train)

# Best model
best_xgb = xgb_grid.best_estimator_
y_pred_xgb = best_xgb.predict(X_test)

# Evaluate
mse_xgb = mean_squared_error(y_test, y_pred_xgb)
r2_xgb = r2_score(y_test, y_pred_xgb)

print(f"Tuned XGBoost - Best Params: {xgb_grid.best_params_}")
```

```
print(f"Tuned XGBoost - MSE: {mse_xgb:.6f}, R2: {r2_xgb:.6f}")
```

```
Fitting 3 folds for each of 108 candidates, totalling 324 fits
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=5,
n_estimators=200, subsample=0.8; total time= 19.7s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=5,
n_estimators=200, subsample=0.8; total time= 19.9s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=5,
n_estimators=200, subsample=0.8; total time= 20.1s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=5,
n_estimators=200, subsample=1.0; total time= 20.0s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=5,
n_estimators=200, subsample=1.0; total time= 20.3s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=5,
n_estimators=200, subsample=1.0; total time= 21.1s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=5,
n_estimators=300, subsample=0.8; total time= 27.2s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=5,
n_estimators=300, subsample=1.0; total time= 27.0s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=5,
n_estimators=300, subsample=0.8; total time= 27.8s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=5,
n_estimators=300, subsample=1.0; total time= 27.9s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=5,
n_estimators=300, subsample=0.8; total time= 28.3s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=5,
n_estimators=300, subsample=1.0; total time= 28.4s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=7,
n_estimators=200, subsample=0.8; total time= 24.8s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=7,
n_estimators=200, subsample=1.0; total time= 25.3s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=7,
n_estimators=200, subsample=0.8; total time= 25.6s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=7,
n_estimators=200, subsample=1.0; total time= 25.6s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=7,
n_estimators=200, subsample=1.0; total time= 26.4s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=7,
n_estimators=200, subsample=0.8; total time= 29.1s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=5,
n_estimators=500, subsample=1.0; total time= 37.7s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=5,
n_estimators=500, subsample=0.8; total time= 38.7s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=5,
n_estimators=500, subsample=0.8; total time= 38.6s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=5,
n_estimators=500, subsample=0.8; total time= 39.2s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=5,
```

```

n_estimators=500, subsample=1.0; total time= 38.8s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=5,
n_estimators=500, subsample=1.0; total time= 40.4s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=7,
n_estimators=300, subsample=0.8; total time= 35.3s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=7,
n_estimators=300, subsample=0.8; total time= 35.6s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=7,
n_estimators=300, subsample=0.8; total time= 36.2s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=7,
n_estimators=300, subsample=1.0; total time= 36.8s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=7,
n_estimators=300, subsample=1.0; total time= 37.8s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=7,
n_estimators=300, subsample=1.0; total time= 36.7s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=7,
n_estimators=500, subsample=0.8; total time= 49.8s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=7,
n_estimators=500, subsample=0.8; total time= 49.8s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=7,
n_estimators=500, subsample=1.0; total time= 50.2s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=7,
n_estimators=500, subsample=0.8; total time= 51.2s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=7,
n_estimators=500, subsample=1.0; total time= 50.6s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=7,
n_estimators=500, subsample=1.0; total time= 51.6s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=9,
n_estimators=200, subsample=0.8; total time= 33.4s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=9,
n_estimators=200, subsample=0.8; total time= 33.6s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=9,
n_estimators=200, subsample=0.8; total time= 33.4s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=9,
n_estimators=200, subsample=1.0; total time= 34.6s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=9,
n_estimators=200, subsample=1.0; total time= 35.0s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=9,
n_estimators=200, subsample=1.0; total time= 34.3s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=9,
n_estimators=300, subsample=0.8; total time= 47.3s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=9,
n_estimators=300, subsample=0.8; total time= 48.0s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=9,
n_estimators=300, subsample=0.8; total time= 47.4s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=9,
n_estimators=300, subsample=1.0; total time= 49.3s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=9,

```

```

n_estimators=300, subsample=1.0; total time= 48.1s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=9,
n_estimators=300, subsample=1.0; total time= 48.6s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=5,
n_estimators=200, subsample=0.8; total time= 15.7s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=5,
n_estimators=200, subsample=0.8; total time= 15.8s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=5,
n_estimators=200, subsample=0.8; total time= 15.8s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=5,
n_estimators=200, subsample=1.0; total time= 15.4s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=5,
n_estimators=200, subsample=1.0; total time= 15.9s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=5,
n_estimators=200, subsample=1.0; total time= 15.3s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=9,
n_estimators=500, subsample=0.8; total time= 1.1min
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=9,
n_estimators=500, subsample=0.8; total time= 1.1min
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=9,
n_estimators=500, subsample=0.8; total time= 1.1min
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=5,
n_estimators=300, subsample=0.8; total time= 20.6s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=5,
n_estimators=300, subsample=0.8; total time= 20.2s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=5,
n_estimators=300, subsample=0.8; total time= 20.4s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=5,
n_estimators=300, subsample=1.0; total time= 20.9s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=5,
n_estimators=300, subsample=1.0; total time= 21.1s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=5,
n_estimators=300, subsample=1.0; total time= 20.8s
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=9,
n_estimators=500, subsample=1.0; total time= 1.2min
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=9,
n_estimators=500, subsample=1.0; total time= 1.2min
[CV] END colsample_bytree=0.8, learning_rate=0.01, max_depth=9,
n_estimators=500, subsample=1.0; total time= 1.2min
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=7,
n_estimators=200, subsample=0.8; total time= 17.3s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=7,
n_estimators=200, subsample=0.8; total time= 18.6s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=7,
n_estimators=200, subsample=1.0; total time= 18.6s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=7,
n_estimators=200, subsample=0.8; total time= 19.1s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=7,

```

```

n_estimators=200, subsample=1.0; total time= 18.5s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=7,
n_estimators=200, subsample=1.0; total time= 19.2s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=5,
n_estimators=500, subsample=0.8; total time= 31.1s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=5,
n_estimators=500, subsample=0.8; total time= 31.4s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=5,
n_estimators=500, subsample=1.0; total time= 30.3s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=5,
n_estimators=500, subsample=0.8; total time= 32.2s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=5,
n_estimators=500, subsample=1.0; total time= 31.6s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=5,
n_estimators=500, subsample=1.0; total time= 38.9s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=7,
n_estimators=300, subsample=0.8; total time= 23.6s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=7,
n_estimators=300, subsample=0.8; total time= 24.5s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=7,
n_estimators=300, subsample=0.8; total time= 24.8s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=7,
n_estimators=300, subsample=1.0; total time= 25.0s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=7,
n_estimators=300, subsample=1.0; total time= 25.4s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=7,
n_estimators=300, subsample=1.0; total time= 26.1s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=7,
n_estimators=500, subsample=0.8; total time= 36.8s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=7,
n_estimators=500, subsample=0.8; total time= 37.7s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=7,
n_estimators=500, subsample=1.0; total time= 36.7s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=9,
n_estimators=200, subsample=0.8; total time= 23.1s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=7,
n_estimators=500, subsample=1.0; total time= 37.6s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=9,
n_estimators=200, subsample=0.8; total time= 23.7s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=7,
n_estimators=500, subsample=0.8; total time= 41.8s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=9,
n_estimators=200, subsample=0.8; total time= 23.0s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=9,
n_estimators=200, subsample=1.0; total time= 23.1s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=9,
n_estimators=200, subsample=1.0; total time= 22.5s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=9,

```

```

n_estimators=200, subsample=1.0; total time= 23.7s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=7,
n_estimators=500, subsample=1.0; total time= 38.0s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=9,
n_estimators=300, subsample=0.8; total time= 28.1s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=9,
n_estimators=300, subsample=0.8; total time= 28.7s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=9,
n_estimators=300, subsample=0.8; total time= 29.1s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=9,
n_estimators=300, subsample=1.0; total time= 29.9s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=9,
n_estimators=300, subsample=1.0; total time= 29.7s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=9,
n_estimators=300, subsample=1.0; total time= 30.1s
[CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=5, n_estimators=200,
subsample=0.8; total time= 14.6s
[CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=5, n_estimators=200,
subsample=0.8; total time= 14.9s
[CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=5, n_estimators=200,
subsample=0.8; total time= 14.7s
[CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=5, n_estimators=200,
subsample=1.0; total time= 14.9s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=9,
n_estimators=500, subsample=1.0; total time= 41.1s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=9,
n_estimators=500, subsample=0.8; total time= 42.7s
[CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=5, n_estimators=200,
subsample=1.0; total time= 14.9s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=9,
n_estimators=500, subsample=0.8; total time= 43.7s
[CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=5, n_estimators=200,
subsample=1.0; total time= 15.3s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=9,
n_estimators=500, subsample=0.8; total time= 42.8s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=9,
n_estimators=500, subsample=1.0; total time= 42.1s
[CV] END colsample_bytree=0.8, learning_rate=0.05, max_depth=9,
n_estimators=500, subsample=1.0; total time= 42.4s
[CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=5, n_estimators=300,
subsample=0.8; total time= 20.0s
[CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=5, n_estimators=300,
subsample=0.8; total time= 20.1s
[CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=5, n_estimators=300,
subsample=0.8; total time= 19.8s
[CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=5, n_estimators=300,
subsample=1.0; total time= 19.8s
[CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=5, n_estimators=300,

```

subsample=1.0; total time= 20.7s
 [CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=5, n_estimators=300,
 subsample=1.0; total time= 20.8s
 [CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=7, n_estimators=200,
 subsample=0.8; total time= 16.7s
 [CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=5, n_estimators=500,
 subsample=0.8; total time= 31.4s
 [CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=5, n_estimators=500,
 subsample=0.8; total time= 31.9s
 [CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=5, n_estimators=500,
 subsample=1.0; total time= 31.7s
 [CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=5, n_estimators=500,
 subsample=1.0; total time= 31.6s
 [CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=5, n_estimators=500,
 subsample=0.8; total time= 31.9s
 [CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=7, n_estimators=200,
 subsample=0.8; total time= 16.3s
 [CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=7, n_estimators=200,
 subsample=0.8; total time= 16.8s
 [CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=7, n_estimators=200,
 subsample=1.0; total time= 16.1s
 [CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=7, n_estimators=200,
 subsample=1.0; total time= 16.5s
 [CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=7, n_estimators=200,
 subsample=1.0; total time= 17.4s
 [CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=5, n_estimators=500,
 subsample=1.0; total time= 31.3s
 [CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=7, n_estimators=300,
 subsample=0.8; total time= 21.8s
 [CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=7, n_estimators=300,
 subsample=0.8; total time= 22.3s
 [CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=7, n_estimators=300,
 subsample=1.0; total time= 22.5s
 [CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=7, n_estimators=300,
 subsample=0.8; total time= 23.1s
 [CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=7, n_estimators=300,
 subsample=1.0; total time= 23.2s
 [CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=7, n_estimators=300,
 subsample=1.0; total time= 24.0s
 [CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=9, n_estimators=200,
 subsample=0.8; total time= 19.6s
 [CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=7, n_estimators=500,
 subsample=0.8; total time= 35.8s
 [CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=7, n_estimators=500,
 subsample=0.8; total time= 36.3s
 [CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=7, n_estimators=500,
 subsample=1.0; total time= 34.9s
 [CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=7, n_estimators=500,

```

subsample=0.8; total time= 35.2s
[CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=7, n_estimators=500,
subsample=1.0; total time= 35.1s
[CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=7, n_estimators=500,
subsample=1.0; total time= 35.6s
[CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=9, n_estimators=200,
subsample=0.8; total time= 18.8s
[CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=9, n_estimators=200,
subsample=0.8; total time= 19.1s
[CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=9, n_estimators=200,
subsample=1.0; total time= 19.3s
[CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=9, n_estimators=200,
subsample=1.0; total time= 19.0s
[CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=9, n_estimators=200,
subsample=1.0; total time= 19.2s
[CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=9, n_estimators=300,
subsample=0.8; total time= 26.1s
[CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=9, n_estimators=300,
subsample=0.8; total time= 26.4s
[CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=9, n_estimators=300,
subsample=1.0; total time= 25.5s
[CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=9, n_estimators=300,
subsample=0.8; total time= 26.9s
[CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=9, n_estimators=300,
subsample=1.0; total time= 26.9s
[CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=9, n_estimators=300,
subsample=1.0; total time= 26.8s
[CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=9, n_estimators=500,
subsample=0.8; total time= 40.9s
[CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=9, n_estimators=500,
subsample=0.8; total time= 40.5s
[CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=9, n_estimators=500,
subsample=0.8; total time= 41.3s
[CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=9, n_estimators=500,
subsample=1.0; total time= 40.9s
[CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=9, n_estimators=500,
subsample=1.0; total time= 40.9s
[CV] END colsample_bytree=0.8, learning_rate=0.1, max_depth=9, n_estimators=500,
subsample=1.0; total time= 41.3s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=5,
n_estimators=200, subsample=0.8; total time= 21.2s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=5,
n_estimators=200, subsample=0.8; total time= 21.6s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=5,
n_estimators=200, subsample=0.8; total time= 21.1s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=5,
n_estimators=200, subsample=1.0; total time= 21.7s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=5,

```



```

n_estimators=200, subsample=1.0; total time= 21.7s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=5,
n_estimators=200, subsample=1.0; total time= 23.1s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=5,
n_estimators=300, subsample=0.8; total time= 29.9s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=5,
n_estimators=300, subsample=0.8; total time= 29.5s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=5,
n_estimators=300, subsample=0.8; total time= 30.1s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=5,
n_estimators=300, subsample=1.0; total time= 31.5s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=5,
n_estimators=300, subsample=1.0; total time= 31.5s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=5,
n_estimators=300, subsample=1.0; total time= 31.3s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=5,
n_estimators=500, subsample=0.8; total time= 42.4s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=5,
n_estimators=500, subsample=0.8; total time= 43.0s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=5,
n_estimators=500, subsample=0.8; total time= 42.8s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=5,
n_estimators=500, subsample=1.0; total time= 43.8s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=5,
n_estimators=500, subsample=1.0; total time= 43.7s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=5,
n_estimators=500, subsample=1.0; total time= 44.7s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=7,
n_estimators=200, subsample=0.8; total time= 27.3s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=7,
n_estimators=200, subsample=0.8; total time= 28.4s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=7,
n_estimators=200, subsample=0.8; total time= 28.1s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=7,
n_estimators=200, subsample=1.0; total time= 27.9s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=7,
n_estimators=200, subsample=1.0; total time= 28.7s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=7,
n_estimators=200, subsample=1.0; total time= 29.1s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=7,
n_estimators=300, subsample=0.8; total time= 37.9s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=7,
n_estimators=300, subsample=0.8; total time= 39.3s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=7,
n_estimators=300, subsample=0.8; total time= 39.2s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=7,
n_estimators=300, subsample=1.0; total time= 39.0s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=7,

```

```

n_estimators=300, subsample=1.0; total time= 39.2s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=7,
n_estimators=300, subsample=1.0; total time= 38.6s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=7,
n_estimators=500, subsample=0.8; total time= 54.6s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=7,
n_estimators=500, subsample=0.8; total time= 55.1s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=7,
n_estimators=500, subsample=0.8; total time= 54.5s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=7,
n_estimators=500, subsample=1.0; total time= 56.6s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=7,
n_estimators=500, subsample=1.0; total time= 59.4s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=9,
n_estimators=200, subsample=0.8; total time= 37.6s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=7,
n_estimators=500, subsample=1.0; total time= 59.0s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=9,
n_estimators=200, subsample=0.8; total time= 37.2s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=9,
n_estimators=200, subsample=0.8; total time= 37.4s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=9,
n_estimators=200, subsample=1.0; total time= 38.7s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=9,
n_estimators=200, subsample=1.0; total time= 38.6s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=9,
n_estimators=200, subsample=1.0; total time= 39.9s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=9,
n_estimators=300, subsample=0.8; total time= 52.1s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=9,
n_estimators=300, subsample=0.8; total time= 53.4s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=9,
n_estimators=300, subsample=0.8; total time= 58.0s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=9,
n_estimators=300, subsample=1.0; total time= 54.3s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=9,
n_estimators=300, subsample=1.0; total time= 55.3s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=9,
n_estimators=300, subsample=1.0; total time= 55.6s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=5,
n_estimators=200, subsample=0.8; total time= 16.9s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=5,
n_estimators=200, subsample=0.8; total time= 16.8s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=5,
n_estimators=200, subsample=0.8; total time= 16.4s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=5,
n_estimators=200, subsample=1.0; total time= 16.6s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=5,

```

```

n_estimators=200, subsample=1.0; total time= 17.0s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=5,
n_estimators=200, subsample=1.0; total time= 16.1s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=9,
n_estimators=500, subsample=0.8; total time= 1.2min
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=9,
n_estimators=500, subsample=0.8; total time= 1.2min
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=9,
n_estimators=500, subsample=0.8; total time= 1.2min
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=5,
n_estimators=300, subsample=0.8; total time= 21.5s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=9,
n_estimators=500, subsample=1.0; total time= 1.3min
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=5,
n_estimators=300, subsample=0.8; total time= 22.2s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=9,
n_estimators=500, subsample=1.0; total time= 1.3min
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=5,
n_estimators=300, subsample=0.8; total time= 21.5s
[CV] END colsample_bytree=1.0, learning_rate=0.01, max_depth=9,
n_estimators=500, subsample=1.0; total time= 1.3min
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=5,
n_estimators=300, subsample=1.0; total time= 22.2s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=5,
n_estimators=300, subsample=1.0; total time= 21.5s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=5,
n_estimators=300, subsample=1.0; total time= 21.8s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=7,
n_estimators=200, subsample=0.8; total time= 18.5s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=7,
n_estimators=200, subsample=0.8; total time= 19.7s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=5,
n_estimators=500, subsample=0.8; total time= 32.6s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=5,
n_estimators=500, subsample=0.8; total time= 33.8s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=7,
n_estimators=200, subsample=0.8; total time= 19.6s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=7,
n_estimators=200, subsample=1.0; total time= 19.3s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=5,
n_estimators=500, subsample=0.8; total time= 35.2s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=7,
n_estimators=200, subsample=1.0; total time= 19.2s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=7,
n_estimators=200, subsample=1.0; total time= 19.2s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=5,
n_estimators=500, subsample=1.0; total time= 33.5s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=5,

```

```

n_estimators=500, subsample=1.0; total time= 32.0s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=5,
n_estimators=500, subsample=1.0; total time= 34.6s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=7,
n_estimators=300, subsample=0.8; total time= 23.9s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=7,
n_estimators=300, subsample=0.8; total time= 25.5s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=7,
n_estimators=300, subsample=1.0; total time= 25.3s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=7,
n_estimators=300, subsample=0.8; total time= 25.7s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=7,
n_estimators=300, subsample=1.0; total time= 26.4s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=7,
n_estimators=300, subsample=1.0; total time= 26.3s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=7,
n_estimators=500, subsample=0.8; total time= 38.3s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=7,
n_estimators=500, subsample=0.8; total time= 38.2s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=7,
n_estimators=500, subsample=0.8; total time= 39.3s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=9,
n_estimators=200, subsample=0.8; total time= 24.2s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=7,
n_estimators=500, subsample=1.0; total time= 37.5s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=7,
n_estimators=500, subsample=1.0; total time= 37.9s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=9,
n_estimators=200, subsample=0.8; total time= 23.9s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=7,
n_estimators=500, subsample=1.0; total time= 37.0s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=9,
n_estimators=200, subsample=0.8; total time= 24.1s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=9,
n_estimators=200, subsample=1.0; total time= 24.9s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=9,
n_estimators=200, subsample=1.0; total time= 24.7s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=9,
n_estimators=200, subsample=1.0; total time= 25.3s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=9,
n_estimators=300, subsample=0.8; total time= 31.0s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=9,
n_estimators=300, subsample=0.8; total time= 30.7s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=9,
n_estimators=300, subsample=0.8; total time= 30.4s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=9,
n_estimators=300, subsample=1.0; total time= 30.5s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=9,

```

```

n_estimators=300, subsample=1.0; total time= 31.5s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=9,
n_estimators=300, subsample=1.0; total time= 31.1s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, n_estimators=200,
subsample=0.8; total time= 15.1s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, n_estimators=200,
subsample=0.8; total time= 15.1s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, n_estimators=200,
subsample=0.8; total time= 15.6s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=9,
n_estimators=500, subsample=0.8; total time= 44.0s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, n_estimators=200,
subsample=1.0; total time= 15.0s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, n_estimators=200,
subsample=1.0; total time= 15.1s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, n_estimators=200,
subsample=1.0; total time= 14.9s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=9,
n_estimators=500, subsample=0.8; total time= 44.6s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=9,
n_estimators=500, subsample=0.8; total time= 44.7s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=9,
n_estimators=500, subsample=1.0; total time= 44.3s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=9,
n_estimators=500, subsample=1.0; total time= 44.3s
[CV] END colsample_bytree=1.0, learning_rate=0.05, max_depth=9,
n_estimators=500, subsample=1.0; total time= 45.0s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, n_estimators=300,
subsample=0.8; total time= 21.0s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, n_estimators=300,
subsample=0.8; total time= 21.2s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, n_estimators=300,
subsample=0.8; total time= 21.3s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, n_estimators=300,
subsample=1.0; total time= 20.6s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, n_estimators=300,
subsample=1.0; total time= 20.7s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, n_estimators=300,
subsample=1.0; total time= 21.1s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, n_estimators=500,
subsample=0.8; total time= 32.4s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, n_estimators=500,
subsample=0.8; total time= 32.6s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, n_estimators=500,
subsample=0.8; total time= 32.3s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=7, n_estimators=200,
subsample=0.8; total time= 16.4s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=7, n_estimators=200,

```

```

subsample=0.8; total time= 17.5s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, n_estimators=500,
subsample=1.0; total time= 32.2s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, n_estimators=500,
subsample=1.0; total time= 33.4s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=7, n_estimators=200,
subsample=0.8; total time= 17.3s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=7, n_estimators=200,
subsample=1.0; total time= 16.7s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=7, n_estimators=200,
subsample=1.0; total time= 16.8s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=5, n_estimators=500,
subsample=1.0; total time= 32.2s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=7, n_estimators=200,
subsample=1.0; total time= 20.3s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=7, n_estimators=300,
subsample=0.8; total time= 22.6s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=7, n_estimators=300,
subsample=0.8; total time= 24.0s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=7, n_estimators=300,
subsample=0.8; total time= 23.1s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=7, n_estimators=300,
subsample=1.0; total time= 23.8s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=7, n_estimators=300,
subsample=1.0; total time= 23.9s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=7, n_estimators=300,
subsample=1.0; total time= 24.7s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=7, n_estimators=500,
subsample=0.8; total time= 35.4s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=7, n_estimators=500,
subsample=1.0; total time= 34.8s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=7, n_estimators=500,
subsample=1.0; total time= 35.6s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=7, n_estimators=500,
subsample=0.8; total time= 36.2s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=7, n_estimators=500,
subsample=0.8; total time= 36.8s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=9, n_estimators=200,
subsample=0.8; total time= 20.0s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=9, n_estimators=200,
subsample=0.8; total time= 19.9s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=7, n_estimators=500,
subsample=1.0; total time= 36.7s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=9, n_estimators=200,
subsample=0.8; total time= 20.3s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=9, n_estimators=200,
subsample=1.0; total time= 19.8s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=9, n_estimators=200,

```

```

subsample=1.0; total time= 19.4s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=9, n_estimators=200,
subsample=1.0; total time= 19.8s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=9, n_estimators=300,
subsample=0.8; total time= 26.9s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=9, n_estimators=300,
subsample=0.8; total time= 26.4s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=9, n_estimators=300,
subsample=1.0; total time= 26.6s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=9, n_estimators=300,
subsample=1.0; total time= 26.7s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=9, n_estimators=300,
subsample=1.0; total time= 25.7s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=9, n_estimators=300,
subsample=0.8; total time= 32.8s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=9, n_estimators=500,
subsample=1.0; total time= 31.0s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=9, n_estimators=500,
subsample=0.8; total time= 31.8s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=9, n_estimators=500,
subsample=0.8; total time= 33.4s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=9, n_estimators=500,
subsample=0.8; total time= 31.6s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=9, n_estimators=500,
subsample=1.0; total time= 30.8s
[CV] END colsample_bytree=1.0, learning_rate=0.1, max_depth=9, n_estimators=500,
subsample=1.0; total time= 30.1s
Tuned XGBoost - Best Params: {'colsample_bytree': 0.8, 'learning_rate': 0.01,
' max_depth': 9, 'n_estimators': 500, 'subsample': 0.8}
Tuned XGBoost - MSE: 0.093542, R2: 0.704931

```

Ensemble (XGBoost + RF)

```

[11]: from sklearn.ensemble import StackingRegressor, RandomForestRegressor,
      ↪ GradientBoostingRegressor
      from sklearn.linear_model import Ridge # Final meta-learner

      # Define base models
      base_models = [
          ('xgb', best_xgb), # Best tuned XGBoost
          ('rf', RandomForestRegressor(n_estimators=200, random_state=42)),
          ('gbr', GradientBoostingRegressor(n_estimators=150, learning_rate=0.05,
      ↪ random_state=42))
      ]

      # Stacking Ensemble
      stacking_model = StackingRegressor(estimators=base_models,
      ↪ final_estimator=Ridge())

```

```

# Train Stacking Model
stacking_model.fit(X_train, y_train)
y_pred_stack = stacking_model.predict(X_test)

# Evaluate
mse_stack = mean_squared_error(y_test, y_pred_stack)
r2_stack = r2_score(y_test, y_pred_stack)

print(f"Stacking Ensemble - MSE: {mse_stack:.6f}, R2: {r2_stack:.6f}")

# ---- PLOT COMPARISON ----
models = ['Tuned XGBoost', 'Stacking Ensemble']
mse_scores = [mse_xgb, mse_stack]
r2_scores = [r2_xgb, r2_stack]

fig, ax = plt.subplots(1, 2, figsize=(12, 5))


# MSE Plot
ax[0].bar(models, mse_scores, color=['blue', 'purple'])
ax[0].set_ylabel("Mean Squared Error (MSE)")
ax[0].set_title("MSE Comparison")

# R2 Score Plot
ax[1].bar(models, r2_scores, color=['blue', 'purple'])
ax[1].set_ylabel("R2 Score")
ax[1].set_title("R2 Comparison")

plt.show()

```

Stacking Ensemble - MSE: 0.092646, R2: 0.707756



Graph Based learning

```
[12]: import networkx as nx
from sklearn.kernel_ridge import KernelRidge
from sklearn.neighbors import kneighbors_graph
from sklearn.metrics import mean_squared_error, r2_score

# ---- CREATE A KNN GRAPH ----
knn_graph = kneighbors_graph(X_train, n_neighbors=5, mode='connectivity',
                             include_self=False)

# Convert to NetworkX Graph
G = nx.from_scipy_sparse_array(knn_graph)

# Visualize Graph Structure
nx.draw(G, node_size=10)

# ---- GRAPH-BASED REDSHIFT ESTIMATION ----
kr = KernelRidge(alpha=1.0, kernel='rbf') # RBF kernel smooths over the graph
kr.fit(X_train, y_train) # Train using photometric data

y_pred_graph = kr.predict(X_test) # Predict redshift

# Evaluate Performance
mse_graph = mean_squared_error(y_test, y_pred_graph)
r2_graph = r2_score(y_test, y_pred_graph)

print(f"Graph-Based Kernel Ridge Regression - MSE: {mse_graph:.6f}, R2: {r2_graph:.6f}")
```

Graph-Based Kernel Ridge Regression - MSE: 0.200561, R2: 0.367346



Hybrid Graph based and XG Boost

```
[13]: # ---- USE GRAPH-BASED REDSHIFT AS A FEATURE ----
df_clean['graph_pred_redshift'] = kr.predict(X) # FIXED: Use Kernel Ridge
      ↪ Regression Predictions

# Retrain XGBoost with Graph-Enhanced Features
X_new = df_clean.drop(columns=['specz_redshift']) # Keep everything except
      ↪ actual redshift
y_new = df_clean['specz_redshift']

X_train_new, X_test_new, y_train_new, y_test_new = train_test_split(X_new,
      ↪ y_new, test_size=0.2, random_state=42)

xgb = XGBRegressor(n_estimators=300, learning_rate=0.05, max_depth=7,
      ↪ random_state=42)
xgb.fit(X_train_new, y_train_new)

y_pred_xgb_new = xgb.predict(X_test_new)
```

```

mse_xgb_new = mean_squared_error(y_test_new, y_pred_xgb_new)
r2_xgb_new = r2_score(y_test_new, y_pred_xgb_new)

print(f"Hybrid Graph + XGBoost Model - MSE: {mse_xgb_new:.6f}, R2: {r2_xgb_new:.6f}")

# ---- PLOT COMPARISON ----
models = ['Stacking Ensemble', 'Graph-Based Model', 'Hybrid Graph + XGBoost']
mse_scores = [mse_stack, mse_graph, mse_xgb_new]
r2_scores = [r2_stack, r2_graph, r2_xgb_new]

fig, ax = plt.subplots(1, 2, figsize=(12, 5))

# MSE Plot
ax[0].bar(models, mse_scores, color=['purple', 'orange', 'blue'])
ax[0].set_ylabel("Mean Squared Error (MSE)")
ax[0].set_title("MSE Comparison")

# R2 Score Plot
ax[1].bar(models, r2_scores, color=['purple', 'orange', 'blue'])
ax[1].set_ylabel("R2 Score")
ax[1].set_title("R2 Comparison")

plt.show()

```

Hybrid Graph + XGBoost Model - MSE: 0.095564, R2: 0.698550

